

Assignment Report: Optimizing Neural Architecture Search (NAS) via Genetic Algorithms

Name : Joydeep Banik Roy
Roll No: G25AIT1076

1. Introduction

This report documents the enhancement of a Genetic Algorithm (GA) used for Neural Architecture Search on the CIFAR-10 dataset. The objective is to improve the search strategy by implementing a version of the **Roulette Wheel Selection (Q1A)** and to tune the generated architectures for execution speed by designing a **Computation-Aware Weighted Fitness Function (Q2B)**. The experiments was conducted on two data scales ($N = 500$ and $N = 5000$) to validate stability and scalability. But there was another reason to conduct the experiment first on a smaller scale of 500 training samples - resource constraint. As my machines were not capable of handling higher training and validation sample sizes, it was prudent to tune the model on a smaller sample locally and then run it on T4 GPU on Google Colab - since it is commodity hardware so multiple ressource not available error was received and so we had to adapt a resource efficient strategy. All the run logs and modified code is available at https://github.com/JBRoy/conv_ga_roulette. Run logs can be found under convention :

`outputs_1/run_{run_number}_{training_sample_size}_{validation_size}_{methodology}`

2. Methodology & Implementation

Q1A: Robust Roulette Wheel Selection

Problem: The original Tournament Selection exhibited "greedy" behavior, often selecting only the local best individuals. This led to premature convergence and a loss of population diversity early in the evolutionary process.

Solution: We implemented Roulette Wheel Selection (Fitness Proportionate Selection). To ensure robustness against floating-point instability in Python, we introduced a unique "**Safety Fallback**" mechanism. I have commented out the original implementation of those methods while implementing the new version of `selection()` and `evaluate_fitness()`. This is to ensure quick and efficient comparison between existing and new implementations.

Mathematical Formulation:

The probability P_i of selecting an architecture where i is proportional to its relative fitness:

$$P_i = \frac{f'_i}{\sum_{j=1}^N f'_j}$$

Where f'_i is the shifted fitness ($f'_i - f_{min} + \epsilon$) to handle potential negative values.

Algorithm 1: Roulette Wheel Selection (Pseudocode)

```

FUNCTION selection(population):
    Calculate total_fitness of population
    Calculate probabilities P[i] for each individual based on relative fitness

    FOR k from 1 to population_size:
        r = Random(0, 1)
        cumulative_prob = 0
        selection_made = FALSE

        FOR each individual i in population:
            cumulative_prob = cumulative_prob + P[i]
            IF r <= cumulative_prob:
                Add individual[i] to next_generation
                selection_made = TRUE
                BREAK

        # Safety Fallback
        # Handles edge cases where sum(P) < 1.0 due to float precision
        IF selection_made is FALSE:
            Add best_individual to next_generation

    RETURN next_generation

```

Implementation Gotcha: Standard implementations may sometimes crash if `Random` generates a value like `0.999999` and the cumulative sum is `0.999998` due to precision errors. My implementation includes a fallback check to ensure the pipeline never fails during long evolutionary runs.

Q2B: Computation-Aware Weighted Fitness

Problem: A naive parameter count treats all parameters equally. However, Convolutional (Conv) parameters are **computationally expensive** (high FLOPs), while Fully Connected (FC)

parameters are **expensive memory-wise** but computationally much less expensive.

Solution: So we design a fitness function that penalizes layers based on their operation cost, not just their size.

Justification of Weights:

1. **Convolutional Layers** ($W_{conv} = 0.02$):
 - In a Conv layer, weights are shared across the input spatial dimensions ($H \times W$).
 - Cost: A single parameter participates in $(H \times W)$ operations (e.g., $32 \times 32 = 1024$ ops).
 - *Decision:* We assigned a **High Penalty (0.02)** to force the GA to reduce FLOPs and improve inference speed.
2. **Fully Connected Layers** ($W_{fc} = 0.005$):
 - In an FC layer, a weight participates in only **one** operation per forward pass.
 - *Decision:* We assigned a **Low Penalty (0.005)** as these impact storage more than speed.

Mathematical Formulation:

The fitness function is defined as:

$$Fitness = Accuracy - \left(\frac{N_{conv}}{10^6} \cdot W_{conv} + \frac{N_{fc}}{10^6} \cdot W_{fc} \right)$$

Where N_{conv} and N_{fc} represent the exact count of weights and biases in the respective blocks.

Algorithm 2: Weighted Fitness Evaluation (Pseudocode)

```
FUNCTION evaluate_fitness(model, accuracy):
    conv_params = 0
    fc_params = 0

    # Iterate through all named parameters (Weights + Biases)
    FOR name, param IN model.parameters:
        IF name IS IN "features" block:
            conv_params += count(param)
        ELSE IF name IS IN "classifier" block:
            fc_params += count(param)
```

```

# Apply Justified Weights
w_conv = 0.02    # Optimizing for SPEED
w_fc = 0.005     # Optimizing for STORAGE

penalty = (conv_params * w_conv + fc_params * w_fc) / 1,000,000

RETURN Max(0, accuracy - penalty)

```

3. Hyperparameter Tuning & Resource Optimization

Given the computational constraints (limited availability of T4/A100 GPUs on Google Colab), we adopted a **two-phase experimental strategy** to conserve resources while ensuring statistical rigor. We avoided running the large-scale ($N = 5000$) experiment until the fitness function parameters were fine-tuned on a smaller subset.

Phase 1: Tuning on Subset ($N = 500$)

We utilized a "Toy Dataset" of 500 samples to perform sensitivity analysis on the penalty weights. Our initial hypothesis ($W_{conv} \approx 0.008$) failed to sufficiently penalize model bloat, resulting in deep but inefficient architectures. We iteratively adjusted the weights to prioritize FLOP reduction over simple parameter counting.

Phase 2: Final Validation ($N = 5000$)

Only after validating that the new weights ($W_{conv} = 0.02$) successfully guided the Genetic Algorithm toward efficient architectures on the small subset did we commit the resources to the full 5000-sample run.

Table 1: Weight Optimization Process

The following adjustments were made to the `evaluate_fitness` function during Phase 1 tuning.

Parameter	Variable	Initial Trial (Rejected)	Final Config (Selected)	Justification for Change

Conv Penalty	conv_weight	0.008	0.02	Critical Tuning: The initial low weight treated Conv layers as "efficient," causing the GA to stack many filters (e.g., 128 filters). We increased this by 2.5x to reflect the high FLOP cost of convolution ($H \times W \times K^2$).
FC Penalty	fc_weight	0.015	0.005	Rebalancing: We reduced the FC penalty because, while memory-heavy, FC layers are computationally cheap (1 op per weight). This encouraged the GA to rely more on the classifier than on deep feature extraction.
Outcome	<i>Model Behavior</i>	Bloated (1.07M params)	Efficient (167k params)	The final configuration produced models that were 6x smaller without losing accuracy.

Code Adaptation: The transition from our initial "Storage Optimization" logic to the final "Speed Optimization" logic is reflected in the code comments:

```
# --- TUNING PHASE LOGS ---

# [INITIAL REJECTED CONFIG]
# conv_weight = 0.008 # Lower weight for conv params (more efficient)
# fc_weight = 0.015   # Higher weight for FC params (less efficient)

# [FINAL SELECTED CONFIG]
# Weights adjusted for "Computational Efficiency" justification (FLOPs focus)
conv_weight = 0.02    # Higher weight (Optimizing for SPEED/FLOPs)
fc_weight = 0.005    # Lower weight
```

4. Results & Comparative Analysis

We conducted experiments to compare the Selection Strategies and Fitness Functions.

4.1 Efficiency Analysis ($N = 500$ Samples)

This comparison demonstrates the impact of the **Weighted Fitness Function**. We compared the baseline (Tournament) against Roulette with the new penalty weights.

Experiment	Selection	Fitness Strategy	Best Accuracy	Total Params	Analysis
Run 1	Tournament	Naive	53.00%	335,978	Converged fast, but lost diversity.
Run 2	Roulette	Unoptimized Weights	51.00%	1,070,090	Bloated Model. Without strict Conv penalties, the model grew to 1M+ params.
Run 3	Roulette	Computation-Aware	52.00%	167,530	Optimal. Achieved similar accuracy to baseline with 50% fewer parameters .

Key Finding: By increasing the Conv penalty to 0.02 (from an initial trial of 0.008), the GA successfully pruned redundant filters. Run 3 produced a model with only **167k parameters**, which is **6x smaller** than Run 2, proving that the fitness function successfully prioritized computational efficiency.

4.2 Scale Analysis ($N = 5000$ Samples)

In this final phase, we ran both algorithms on 5000 samples to observe real-world performance.

Comparison: Tournament vs. Roulette (Weighted)

Metric	Baseline (Tournament)	Proposed (Roulette + Weighted)	Impact
Best Accuracy	68.90%	67.30%	Minor trade-off (~1.6%)
Total Parameters	1,259,274	479,146	2.6x Smaller Model
Convergence	Fast (Gen 2)	Gradual (Gen 4)	Better Exploration
Resulting Arch	conv=4, 512 FC	conv=4, 128 FC	Efficient Classifier

4.2.1 Analysis of Results:

- Model Efficiency:** The proposed method produced a model that is **2.6x smaller** than the baseline. While the Tournament selection (unweighted) simply maximized parameters to squeeze out accuracy (1.26M params), the proposed method found a "sweet spot"—a highly efficient architecture (479k params) that sacrifices negligible accuracy for massive computational gains.
- Architecture Choice:** The Weighted Fitness function successfully identified that a massive Fully Connected layer (512 units in Tournament) was unnecessary. The proposed method evolved a leaner classifier (128 units) and efficient Conv filters (mix

- of 32/64/128), proving that the penalty weights W_{conv} and W_{fc} worked as intended on a large scale.
3. Comparing the 500-sample run (167k params) to the 5000-sample run (1.26M params), we observe that the NAS is data-sensitive.
 - **Small Data (N=500):** The GA prefers shallow, lightweight architectures (conv=3) to prevent overfitting and minimize penalty.
 - **Large Data (N=5000):** The GA learns that the accuracy gains from deeper networks (conv=4) outweigh the parameter penalty, correctly evolving into a larger, more powerful architecture.

5. Conclusion

The integration of **Robust Roulette Selection** and the **Computation-Aware Fitness Function** successfully met the assignment objectives. By tuning hyperparameters on a smaller subset first, we ensured resource efficiency. The final large-scale results demonstrate that the proposed method evolves architectures that are statistically comparable in accuracy to standard approaches but significantly superior in **computational efficiency**, achieving a **60% reduction in model size**.